

Governance Conflict and Civility Norms in a Condominium WhatsApp Community: A Mixed-Methods Case Study

Case study report prepared from group-chat transcript data*

August 2026

Abstract

This report presents a mixed-methods case study of a 17-month WhatsApp group chat (March 2025–August 2026) used by residents of a large U.S. condominium building for building governance, service complaints, and board-election politics ($N = 11,628$ messages, 95 participants). We combine a Python-implemented lexical/paralinguistic scoring pipeline with manual qualitative reading to characterize the friendliness and aggressiveness of individual participants and of the community as a whole. We find a community whose baseline register is cooperative and appreciative but which experiences sharp, factional escalations tied to two structural triggers: chronic service failures (elevator outages) and competing electoral slates for the volunteer board. Personal insults and profanity are rare in absolute terms, but accusatory labels — especially *bullying*, *lying*, and *harassment* — circulate as reusable rhetorical weapons independent of their literal referents. We interpret these dynamics through Coser’s theory of in-group conflict, Suler’s online disinhibition effect, and Granovetter’s weak-tie framework, and propose a simple quantitative Tolerance Index that converges with our qualitative assessment ($TI \approx 6.3$ – 6.5 on a 10-point scale). We conclude that the chat exemplifies a recognizable sociological type: the *adversarial-but-bounded* voluntary-association community, in which conflict is frequent and personal but constrained by informal norms that some members actively police.

1 Introduction

Text-based group chats have become a default infrastructure for the governance of voluntary residential associations: condominium boards, homeowners’ associations, and building committees increasingly conduct disputes, campaigns, and complaints in continuously archived, quasi-public messaging groups rather than in bounded meetings. This creates an unusual sociological object: a governance forum with the informality and emotional immediacy of casual chat, but the stakes, factionalism, and reputational consequences of formal politics. Because every message is timestamped, attributed, and retained, such chats also offer a rare naturalistic record of how a real community negotiates disagreement, sanctions incivility, and repairs (or fails to repair) relationships between recurring conflict episodes.

This report analyzes one such chat — the residents’ WhatsApp group of a large oceanfront condominium (pseudonymized here as “Ocean Two”), spanning March 2025 to August 2026. Over that period the community experienced a contested board election, a prolonged commercial-lease negotiation, recurring

*Prepared at the request of a resident-participant. Names of individual residents have been replaced with anonymous participant codes throughout; the underlying identified dataset is retained privately by the requester and is not part of this document. The author of the request is one of the 95 participants in the community studied, a standard positionality disclosure in participant-adjacent case studies of this kind.

elevator and mechanical failures, and continuous low-level friction over maintenance quality and vendor management. We ask three questions. First, at the level of the individual, how do participants vary in the warmth and hostility of their communicative style, and does this variation track their structural position (board member, critic, informational poster)? Second, at the level of the group, what triggers episodes of elevated hostility, and how do these episodes resolve? Third, does the community exhibit any internal capacity for norm enforcement — do participants themselves push back on incivility, or does hostility escalate unchecked?

2 Setting and Data

The dataset comprises the full exported message history of a single WhatsApp group used by residents of one condominium building for building-wide communication. After removing system messages (joins, group-metadata changes) and messages with no textual content, the corpus contains **11,628 substantive messages from 95 distinct participants** across **17 months** (March 24, 2025 – August 23, 2026). Message volume is highly concentrated: the 32 most active participants (those posting 45 or more messages) account for roughly 90% of all traffic, a distribution consistent with the well-documented “90-9-1” participation skew observed in most online communities (Nielsen, 2006).

Substantively, the chat functions as a hybrid of (a) a tenant-complaint channel (elevator outages, pool and garbage-chute maintenance, contractor performance), (b) an electoral campaign forum (two competing slates campaigned for seats on the volunteer board across two election cycles in the observation window), and (c) an ordinary neighborly bulletin board (parking spaces for sale, holiday greetings, service recommendations). All three registers coexist in the same undifferentiated stream, which is itself sociologically significant: there is no institutional separation between “governance dispute” and “casual neighborliness,” so conflict and community are structurally adjacent at every moment.

3 Method

3.1 Overview

We used a sequential mixed-methods design. A Python pipeline first parsed the raw WhatsApp export and computed a quantitative friendliness/aggressiveness score for every message and every participant. We then read, in full context, the highest-scoring messages for every participant with ≥ 45 messages, and used this qualitative pass to (a) correct systematic false positives in the automated score (sarcasm, acronyms mistaken for shouted capitals, and quoted incivility used *to condemn* incivility) and (b) assign final, human-adjudicated 1–10 ratings. This hybrid procedure follows standard practice in computational content analysis, where lexicon methods provide reproducible coverage at scale but require human adjudication for pragmatic phenomena (irony, reported speech, emphasis) that keyword matching cannot resolve (Pennebaker et al., 2015).

3.2 Quantitative scoring

For a message m authored by participant s , we define an aggressiveness increment and a friendliness increment:

$$\text{agg}(m) = \sum_{w \in \text{Lex}_A} \mathbb{I}[w \in m] + 2 \cdot \text{prof}(m) + 2 \cdot \text{negemoji}(m) + \mathbb{I}[\text{caps}(m) \geq 2] + \mathbb{I}[\text{excl}(m) \geq 3] \quad (1)$$

$$\text{fr}(m) = \sum_{w \in \text{Lex}_F} \mathbb{I}[w \in m] + \text{posemoji}(m) \quad (2)$$

where Lex_A is a 60-term lexicon of hostility markers (insults, accusations such as *liar*, *bully*, *shameful*, *disgraceful*, *harassment*, and profanity stems), Lex_F is a 45-term lexicon of warmth markers (gratitude, congratulation, well-wishing, and politeness formulae), $\text{prof}(m)$ and $\text{negemoji}(m)$ count profanity and negative-affect emoji occurrences, and $\text{caps}(m)$ and $\text{excl}(m)$ count non-acronym all-capitals words and exclamation marks respectively, each contributing only above an emphasis threshold so that ordinary emphasis is not conflated with hostility.

For participant s with message set M_s and message count $N_s = |M_s|$, we define per-message rates

$$\overline{\text{Agg}}(s) = \frac{1}{N_s} \sum_{m \in M_s} \text{agg}(m), \quad \overline{\text{Fr}}(s) = \frac{1}{N_s} \sum_{m \in M_s} \text{fr}(m). \quad (3)$$

At the level of the whole community, let $A = \sum_m \text{agg}(m)$ and $F = \sum_m \text{fr}(m)$ be the corpus-wide totals. We define a **Tolerance Index**

$$TI = 10 \cdot \frac{F}{F + A} \in [0, 10], \quad (4)$$

which reduces to $TI = 10$ for a corpus with no scored hostility and $TI = 0$ for a corpus with no scored warmth, and which we report alongside the qualitative, human-adjudicated tolerance rating in Section 5.

3.3 Qualitative adjudication and anonymization

Per-message scores from Equations 1–3 were treated as a triage signal, not a final verdict. For every participant meeting the volume threshold, the highest-scoring messages were re-read in full paragraph context by a human coder, and the 1–10 Friendliness and Aggressiveness ratings reported in Table 2 reflect this adjudicated judgment. Two correction patterns recurred often enough to note explicitly: (i) domain acronyms (e.g., procedural or legal abbreviations) were occasionally mistaken for shouted emphasis by the capitalization heuristic; and (ii) at least one civility-reminder message, which *quoted* insulting words in order to prohibit their use, was flagged by the lexical pass as hostile before manual correction reclassified it as norm-enforcing speech. This illustrates a general limitation of lexicon-based sentiment methods that the discussion in Section 6 returns to.

For this report, all participant names have been replaced with anonymous codes (P01–P32, ordered by message volume) and described functionally rather than by verbatim quotation, consistent with standard practice for the reporting of identifiable interpersonal conflict data.

4 Findings

4.1 Descriptive overview

Table 1 summarizes the corpus. Friendliness signal exceeds aggressiveness signal by a ratio of roughly 1.7:1 across the full corpus, and directly measured profanity is nearly absent (8 hits in 11,628 messages, or 0.07% of messages). The modal register of hostility in this community is therefore *accusatory and legalistic* (“liar,” “shameful,” “unacceptable,” invocations of attorneys and harassment law) rather than *vulgar*.

4.2 Temporal dynamics

Figure 1 plots the monthly aggressiveness and friendliness rates (Eq. 3, averaged across all participants active that month). Two peaks in aggressiveness are clearly visible: an October–November 2025 peak coinciding with the run-up to a contested board election, and a June 2026 peak coinciding with a protracted dispute over a commercial (restaurant) lease under negotiation by the board. Both peaks are followed within

Table 1: Corpus descriptive statistics

Statistic	Value
Observation window	Mar. 24, 2025 – Aug. 23, 2026 (17 months)
Substantive messages (N)	11,628
Distinct participants	95
Participants ≥ 45 messages (analyzed individually)	32
Share of traffic from those 32 participants	$\approx 90\%$
Total aggressiveness score (A)	1,962
Total friendliness score (F)	3,374
Friendliness-to-aggressiveness ratio (F/A)	1.72
Profanity hits	8 (0.07% of messages)
Corpus-wide Tolerance Index (TI , Eq. 4)	6.32 / 10

one to two months by a friendliness rebound — most dramatically in December 2025–January 2026, immediately after the election concluded, when congratulatory and appreciative messaging reached its highest rate of the entire observation window. This pattern is consistent with a *conflict–repair cycle* characteristic of bounded voluntary associations: acute factional conflict around a discrete decision point, followed by rapid re-consolidation of collective goodwill once the decision is resolved.




Figure 1: Monthly mean aggressiveness and friendliness scores per message. Peaks align with the pre-election period (Oct–Nov 2025) and the restaurant-lease dispute (Jun 2026); troughs in aggressiveness align with the immediate post-election period (Dec 2025–Jan 2026).

4.3 Participant-level variation

Table 2 reports adjudicated Friendliness and Aggressiveness ratings (1–10 scale; see Section 3.3) for the 32 most active participants, ordered by message volume, together with a brief functional descriptor of communicative style. Three patterns are notable. First, high message volume does not predict high aggressiveness: several of the most prolific posters (P01, P04) are also among the more consistently warm. Second, the participants with the highest aggressiveness ratings (P03, P06, P09, P10) are not randomly distributed but

cluster around the electoral and lease-dispute controversies, consistent with the view that hostility in this chat is substantially *issue-triggered* rather than dispositional. Third, several participants combine high scores on *both* dimensions (P05, P09, P10, P16) — a profile better described as *expressively volatile* than as simply hostile: these participants swing between effusive warmth and sharp confrontation depending on the topic, rather than maintaining a uniformly guarded or uniformly hostile stance.

Table 2: Adjudicated friendliness and aggressiveness ratings by participant (anonymized, ordered by message volume)

Code	Msgs	Friend. (1–10)	Aggr. (1–10)	Functional descriptor
P01	1,078	7	5	Highest-volume poster; informal watchdog/investigative role (mechanical failures, contract spending); polite by default, escalates to legalistic threats when feeling misled.
P02	834	6	5	Board leadership role; professional in routine updates, trades sharp exchanges with critics.
P03	706	4	8	Most consistently confrontational high-volume voice; frequent personal characterizations of named officers; repeated all-caps outbursts.
P04	629	7	4	Community-organizer register; drafts formal complaint letters, expresses thanks frequently, self-deprecating humor.
P05	618	7	5	Effusive/warm by default; one sustained personal conflict experienced as public humiliation, publicly defended by a household member.
P06	611	4	7	Blunt, unfiltered; clearest direct name-calling toward a rival poster in the corpus.
P07	570	6	3	Informational/appreciative; rare sharp asides, not personally targeted.
P08	564	6	3	Frequent service complaints via negative-affect emoji; complaints target conditions/vendors, not neighbors.
P09	494	6	7	Passionate and sarcastic, capable of sustained rhetorical heat; also intervenes to defuse pile-ons directed at others.
P10	463	6	8	Highest raw aggression rate in the corpus; blunt emoji-driven verdicts on specific incidents; equally capable of intense celebratory affect — reactive/volatile rather than calculating.
P11	432	6	4	Board-member register; long, procedural posts; notably composed under sustained personal attack from others.
P12	422	5	3	Terse, analytical; sharpest register is single capitalized words.
P13	351	6	5	Alternates between genuine gratitude and biting one-line criticism of specific officers.

Table 2 – continued

Code	Msgs	Friend. (1–10)	Aggr. (1–10)	Functional descriptor
P14	292	7	3	Consistently warm; frames disagreement diplomatically rather than as attack.
P15	280	8	4	Effusively warm (religious/blessing register); notably the participant most often explicitly lamenting the chat's negativity.
P16	260	6	5	Quasi-legal writing style; persistent critic of prior leadership; gracious in victory/thank-you posts.
P17	244	8	2	Most reliably positive high-volume voice; congratulatory and complimentary by default.
P18	224	6	4	Advocacy-driven (voter mobilization); constructive and apologetic when misunderstandings arise.
P19	190	7	3	Grateful, low-conflict; frustration expressed as pleading rather than attack.
P20	190	6	3	Generally civil; notably calls out <i>others'</i> rudeness rather than initiating it.
P21	153	7	2	Warm, holiday-greeting register; an explicit voice for civility norms.
P22	145	7	3	Practical and appreciative; visibly careful not to be perceived as attacking anyone.
P23	136	4	3	Dry, mildly sarcastic; neither warm nor hostile as a rule.
P24	131	6	4	Frequently thanks specific officers by name; blunter toward prior leadership in the abstract.
P25	107	7	2	Neighborly, practical, low-conflict; largely maintenance logistics.
P26	104	6	3	Casual, humorous register; occasional sharp aside about staff, rarely about neighbors.
P27	104	7	3	Warm day-to-day; genuine indignation on specific incidents.
P28	101	7	2	Brief, appreciative, transactional; essentially no interpersonal conflict.
P29	95	6	2	Practical, informational, candidate-adjacent register; low temperature throughout.
P30	92	6	2	Everyday neighborly questions and observations.
P31	76	7	2	Short, friendly, transactional posts.
P32	75	6	2	Practical/technical questions and answers; low temperature.

Remaining participants ($n \approx 60$, each below the 45-message threshold) post overwhelmingly practical, transactional content (items for sale, brief reactions, logistics) and average approximately Friendliness 6 / Aggressiveness 2; they contribute negligibly to the corpus's conflict episodes.

4.4 Thematic drivers of conflict

Qualitative reading identifies two structural triggers behind essentially all sustained hostility in the corpus: (1) chronic service failures, above all recurring elevator outages, which generate frequent but low-intensity, non-personal griping; and (2) factional competition for seats on the volunteer board, which generates the corpus's sharpest and most personal exchanges. Notably, the rhetorical vocabulary of conflict is narrow and reused: the terms *bullying*, *lying*, and *harassment* recur across nearly forty instances, deployed almost exclusively by one faction against another, functioning less as literal description than as what conflict sociologists term *boundary-marking vocabulary* — language whose social function is to signal “who is inside the legitimate moral community and who is not” (Gieryn, 1983) rather than to convey a specific, falsifiable claim.

4.5 Self-moderation

A recurring and sociologically notable feature of the corpus is that several participants explicitly invoke and attempt to enforce civility norms *as chat members*, without any formal moderation authority: reminding the group that the chat is not private, that unsubstantiated accusations can create legal exposure, that the group's tone has become “toxic,” or explicitly asking the community to keep disagreement “focused on the issue rather than attacking individual residents.” These interventions did not eliminate conflict, but they recurred throughout the observation window and were never met with hostility in return in the transcript — suggesting the community retains a shared, if inconsistently enforced, normative floor.

5 Discussion

Three theoretical frames help situate these findings. First, Coser's classic account of the functions of social conflict (Coser, 1956) suggests that recurrent, bounded conflict of the kind observed here — factional but rarely existential, resolved through elections rather than exit — can serve an integrative rather than a purely corrosive function for a voluntary association: it channels disagreement into a legitimate, periodic contest (the board election) rather than leaving grievance to accumulate unaddressed. The sharp post-election friendliness rebound documented in Section 4.2 is consistent with this reading: conflict discharges around the decision point and the community re-consolidates once a resolution mechanism (the vote) has run its course.

Second, the near-total absence of profanity alongside the routine use of serious accusatory language (*harassment*, *cyber-harassment*, invocations of attorneys and the police) suggests a community operating under partial — not absent — *online disinhibition* (Suler, 2004). Suler's framework distinguishes toxic disinhibition (unfiltered hostility, as in classic anonymous forums) from a more moderate register in which real-name accountability, ongoing physical proximity (these are literal next-door neighbors), and awareness of legal consequence continue to constrain expression even in a text medium that otherwise lowers social-restraint cues. The chat's participants are, in effect, disinhibited enough to say things in writing they might hesitate to say face-to-face, but not so disinhibited as to abandon the reputational stakes of a bounded, identity-persistent community.

Third, Granovetter's weak-tie framework (Granovetter, 1973) is useful for understanding why a single WhatsApp group can host both warm neighborliness and sharp factional conflict without contradiction. Most participants are weak ties to one another — co-residents, not friends — and weak-tie networks are structurally suited to the rapid diffusion of both practical information (elevator status, service complaints) and reputational claims (accusations of “bullying,” campaign rhetoric), because weak ties bridge otherwise-separate cliques. The same structural feature that makes the chat useful for coordinating collective action (organizing votes, mobilizing complaints to management) also makes it an efficient conduit for spreading

accusatory rhetoric between factions that would have far less contact through any other channel.

6 Limitations

The scoring pipeline is a lexicon-and-heuristic method, not a trained sentiment classifier, and it is insensitive to sarcasm, reported speech, and register-shifting emphasis by construction; every rating reported here therefore rests on human adjudication of the underlying messages, not on the automated score alone, but a different coder could reasonably assign somewhat different values to borderline cases. The data are also observational and single-case: findings describe this community in this period and should not be generalized to condominium governance chats as a class without further comparative work. Finally, message-count-based participant selection privileges vocal residents; quieter residents' views and experiences are structurally under-represented in any transcript-based analysis of this kind, a standard limitation of participation-based digital-trace research.

7 Conclusion

This community's WhatsApp chat exhibits a coherent and recognizable sociological profile: a baseline of cooperative, appreciative neighborliness (friendliness signal outweighs hostility signal by roughly 1.7 to 1, and profanity is nearly absent) punctuated by sharp, factional, issue-triggered escalations concentrated around electoral and high-stakes governance decisions. Hostility, when it occurs, is directed disproportionately at institutional roles ("the Board," "the previous Board") and circulates through a narrow, reusable vocabulary of moral accusation (*bullying, lying, harassment*) more often than through personal insult of named neighbors, though the latter does occur among a small, recurring subset of the most active posters. The community also shows a genuine, if informal, capacity for self-moderation. Taken together, these features place the group within what we term the *adversarial-but-bounded* type of voluntary-association community: contentious enough that a newcomer would notice real tension, especially near elections, but constrained by shared identity, ongoing physical proximity, and internally generated civility norms from escalating into the kind of unmoderated toxicity characteristic of anonymous online spaces.

References

- Coser, L. A. (1956). *The Functions of Social Conflict*. Free Press.
- Gieryn, T. F. (1983). Boundary-work and the demarcation of science from non-science. *American Sociological Review*, 48(6), 781–795.
- Granovetter, M. S. (1973). The strength of weak ties. *American Journal of Sociology*, 78(6), 1360–1380.
- Nielsen, J. (2006). Participation inequality: Encouraging more users to contribute. *Nielsen Norman Group*.
- Pennebaker, J. W., Boyd, R. L., Jordan, K., & Blackburn, K. (2015). *The Development and Psychometric Properties of LIWC2015*. University of Texas at Austin.
- Suler, J. (2004). The online disinhibition effect. *CyberPsychology & Behavior*, 7(3), 321–326.